

STT LAB Assignment - 11

[Repo](#)

Name	Roll number
<i>Parthiv Patel</i>	23110237
<i>Aryan Solanki</i>	23110049

```
▶ ▾  
# Load .tsv files  
train_df = pd.read_csv("train.tsv", sep='\t')  
test_df = pd.read_csv("test.tsv", sep='\t')  
  
train_data, val_data = train_test_split(train_df, test_size=0.2, random_state=42)  
  
# Check results  
print("Training set:", train_data.shape)  
print("Validation set:", val_data.shape)  
[25]  
... Training set: (5535, 2)  
Validation set: (1384, 2)  
Python
```

```
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split

# Load with column names manually set
train_df = pd.read_csv("train.tsv", sep='\t', header=None, names=["sentence", "label"])
test_df = pd.read_csv("test.tsv", sep='\t', header=None, names=["sentence", "label"])

# Split 20% for validation
train_data, val_data = train_test_split(train_df, test_size=0.2, random_state=42)

# Initialize Bag-of-Words vectorizer
vectorizer = CountVectorizer(max_features=10000)
X_train = vectorizer.fit_transform(train_data['sentence']).toarray()
X_val = vectorizer.transform(val_data['sentence']).toarray()
y_train = train_data['label'].values
y_val = val_data['label'].values

print("X_train shape:", X_train.shape)
print("X_val shape:", X_val.shape)
```

[26] Python

```
... X_train shape: (5536, 10000)
X_val shape: (1384, 10000)
```

```
import torch.nn as nn
import torch

class MLP(nn.Module):
    def __init__(self, input_dim, hidden_1=512, hidden_2=256, hidden_3=128, hidden_4=64, output_dim=2):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(input_dim, hidden_1)
        self.fc2 = nn.Linear(hidden_1, hidden_2)
        self.fc3 = nn.Linear(hidden_2, hidden_3)
        self.fc4 = nn.Linear(hidden_3, hidden_4)
        self.fc5 = nn.Linear(hidden_4, output_dim)
        self.relu = nn.ReLU()
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc2(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc3(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc4(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc5(x)
        return x
```

Python

```

model = MLP(input_dim=10000)

# Count trainable parameters
num_params = sum(p.numel() for p in model.parameters() if p.requires_grad)

print(f"Total trainable parameters: {num_params:,}")

```

```

"""
(10000 * 512 + 512) = 5,120,512
(512 * 256 + 256) = 131,328
(256 * 128 + 128) = 32,896
(128 * 64 + 64) = 8,256
(64 * 2 + 2) = 130
-----
Total = 5,293,122

```

```

# 5,293,122 trainable parameters
"""

```

Python

Total trainable parameters: 5,293,122

```

train_losses, val_losses, val_accuracies = [], [], []

```

```

epoch in range(10):
    model.train()
    running_loss = 0.0
    for batch_x, batch_y in train_loader:
        optimizer.zero_grad()
        outputs = model(batch_x)
        loss = criterion(outputs, batch_y)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    avg_train_loss = running_loss / len(train_loader)
    train_losses.append(avg_train_loss)

```

```

# Validation
model.eval()
val_loss = 0.0
correct = 0
total = 0
with torch.no_grad():
    for batch_x, batch_y in val_loader:
        outputs = model(batch_x)
        loss = criterion(outputs, batch_y)
        val_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct += (predicted == batch_y).sum().item()
        total += batch_y.size(0)
acc = 100 * correct / total
val_losses.append(val_loss / len(val_loader))
val_accuracies.append(acc)

```

```

print(f"Epoch {epoch+1} | Train Loss: {avg_train_loss:.4f} | Val Loss: {val_loss:.4f} | Val Acc: {acc:.2f}%")


```

```

# Save best model
if acc > best_val_acc:
    best_val_acc = acc
    torch.save(model.state_dict(), "checkpoint.pt")

```

Epoch	Train Loss	Val Loss	Val Acc
Epoch 1	0.6082	10.3604	78.68%
Epoch 2	0.2775	10.0236	80.71%
Epoch 3	0.0827	15.5300	79.70%
Epoch 4	0.0207	23.8182	78.68%
Epoch 5	0.0076	29.5870	80.06%
Epoch 6	0.0056	31.8124	79.12%
Epoch 7	0.0028	38.6993	78.97%
Epoch 8	0.0006	42.6305	79.12%
Epoch 9	0.0005	44.7238	79.91%
Epoch 10	0.0001	47.7574	80.20%



Original Model Results:
 Accuracy: 80.56%
 Model Size: 20.20 MB
 Inference Time: 58.56 ms
 Number of parameters: 5,293,122

Dynamic Quantization Results:
 Accuracy: 80.49%
 Model Size: 5.06 MB
 Inference Time: 94.78 ms
 Total trainable parameters (dynamic): 0

Half Precision Results:
 Accuracy: 80.56%
 Model Size: 10.10 MB
 Inference Time: 14319.69 ms
 Total trainable parameters (FP16): 5,293,122

S.I.	Model Name	Test Accuracy (Out of 100)	Storage (In MB)	Number of parameters using in-build function	Inference time (In ms)
1.	Original	80.56%	20.20 MB	5,293,122	58.56 ms
2.	Dynamic	80.49%	5.06 MB	0	94.78 ms
3.	Half	80.56%	10.10 MB	5,293,122	14319.69 ms

Time for Half is more as it lacks support on CPU.